

# COMP6210 Automated Software Verification

## Automata-Based Verification for LTL

---

Pavel Naumov

# Intended Learning Outcomes

By the end of this lecture, you will be able to

- define a Büchi automaton
- construct Buchi automata for simple LTL formulas
- explain how Buchi automata are used in model-checking LTL formulas


Büchi Automata

LTL Model Checking

# Recap on Finite Automata (1)

## different btw fa and ts

- **finite automata** accept/reject *strings* over a given *alphabet*
- equivalent to regular expressions
- a finite automaton over the alphabet  $\{a, b\}$ :


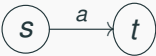




- the initial state is  $s_1$
- the final state is  $s_1$
- the string **abba** is accepted
- the string **bab** is rejected

## Recap on Finite Automata (2)

Let  $\Sigma$  be a finite **alphabet** (consisting symbols that we will use in strings/words over  $\Sigma$ ).

A **finite automaton**  $\mathbb{A}$  consists of:

- a finite set of **states**: 
- a finite set of **transitions** between states, labelled by symbols from  $\Sigma$ : 
- a subset of **initial states**: 
- a subset of **final/accepting states**: 

## Recap on Finite Automata (3)

- A **run** of an automaton on a finite word  $a_1 \dots a_n$  is a finite path through the automaton starting in an initial state, with transitions labelled by  $a_1, \dots, a_n$  (in this order):

$$s_0 \xrightarrow{a_1} s_1 \xrightarrow{a_2} \dots \xrightarrow{a_{n-1}} s_{n-1} \xrightarrow{a_n} s_n$$

- An automaton  $\mathbb{A}$  **accepts** a finite word  $a_1 \dots a_n$  if there exists a run of  $\mathbb{A}$  on  $a_1, \dots, a_n$  which ends in an accepting state.
- The **language** of an automaton  $\mathbb{A}$ , denoted  $L(\mathbb{A})$ , consists of all the words it accepts.


## Recap on Finite Automata (3)

- A **run** of an automaton on a finite word  $a_1 \dots a_n$  is a finite path through the automaton starting in an initial state, with transitions labelled by  $a_1, \dots, a_n$  (in this order):

$$s_0 \xrightarrow{a_1} s_1 \xrightarrow{a_2} \dots \xrightarrow{a_{n-1}} s_{n-1} \xrightarrow{a_n} s_n$$

- An automaton  $\mathbb{A}$  **accepts** a finite word  $a_1 \dots a_n$  if there **exists** a run of  $\mathbb{A}$  on  $a_1, \dots, a_n$  which ends in an accepting state.
- The **language** of an automaton  $\mathbb{A}$ , denoted  $L(\mathbb{A})$ , consists of all the words it accepts.

## Recap on Finite Automata (4)



- some runs of this automaton:

$$s_1 \xrightarrow{b} s_2 \xrightarrow{b} s_2 \xrightarrow{a} s_1$$

$$s_1 \xrightarrow{a} s_1 \xrightarrow{b} s_2 \xrightarrow{b} s_2 \xrightarrow{a} s_1$$


- some words accepted by this automaton:

$\epsilon$  (the empty word),  $a$ ,  $ba$ ,  $bba$ ,  $baba$ , ...

- the language of this automaton consists of:
  - the empty word,
  - all non-empty words that end with an  $a$ .



# Finite Automata over Infinite Words (Büchi Automata)




(The initial state is  $s_1$ .)

- An **infinite run** is an infinite path through the automaton, starting in an initial state.

- e.g.  $s_0 \xrightarrow{a_1} s_1 \xrightarrow{a_2} s_2 \xrightarrow{a_3} \dots$

- An automaton **accepts** an infinite word  $a_1 a_2 a_3 \dots$  if there exists a run of the automaton labelled by  $a_1, a_2, \dots$ , which passes through an accepting state **infinitely often**. We call such a run **accepting**.

# Finite Automata over Infinite Words (Büchi Automata)




(The initial state is  $s_1$ .)

- An **infinite run** is an infinite path through the automaton, starting in an initial state.

- e.g.  $s_0 \xrightarrow{a_1} s_1 \xrightarrow{a_2} s_2 \xrightarrow{a_3} \dots$

- An automaton **accepts** an infinite word  $a_1 a_2 a_3 \dots$  if there exists a run of the automaton labelled by  $a_1, a_2, \dots$ , which passes through an accepting state **infinitely often**. We call such a run **accepting**.

## Example




(The only initial state and only final state is  $s_1$ .)

The run  $s_1 \xrightarrow{b} s_2 \xrightarrow{a} s_1 \xrightarrow{b} s_2 \xrightarrow{a} \dots$  is accepting.

The run  $s_1 \xrightarrow{a} s_1 \xrightarrow{b} s_2 \xrightarrow{b} s_2 \xrightarrow{b} \dots$  is not accepting.

The language of this Büchi automaton consists of all **infinite** words containing infinitely many  $a$  symbols.

## Example



(The only initial state and only final state is  $s_1$ .)

The run  $s_1 \xrightarrow{b} s_2 \xrightarrow{a} s_1 \xrightarrow{b} s_2 \xrightarrow{a} \dots$  is accepting.

The run  $s_1 \xrightarrow{a} s_1 \xrightarrow{b} s_2 \xrightarrow{b} s_2 \xrightarrow{b} \dots$  is not accepting.

The language of this Büchi automaton consists of all **infinite** words containing infinitely many  $a$  symbols.

## Some Properties of Büchi Automata

If  $\mathbb{A}$  and  $\mathbb{B}$  are automata, then there exist:


- an automaton  $\mathbb{A} \cap \mathbb{B}$  (the **product** automaton) such that
$$L(\mathbb{A} \cap \mathbb{B}) = L(\mathbb{A}) \cap L(\mathbb{B}).$$
- an automaton  $\overline{\mathbb{A}}$  (the **complement** automaton) such that  $L(\overline{\mathbb{A}})$  contains exactly those infinite words which are **not** in  $L(\mathbb{A})$ .

## Some Properties of Büchi Automata

If  $\mathbb{A}$  and  $\mathbb{B}$  are automata, then there exist:


- an automaton  $\mathbb{A} \cap \mathbb{B}$  (the **product** automaton) such that
$$L(\mathbb{A} \cap \mathbb{B}) = L(\mathbb{A}) \cap L(\mathbb{B}).$$
- an automaton  $\overline{\mathbb{A}}$  (the **complement** automaton) such that  $L(\overline{\mathbb{A}})$  contains exactly those infinite words which are **not** in  $L(\mathbb{A})$ .

# Buchi Automaton Complement Construction



accepts infinite words that have infinitely many symbols a

how to change that ?



无限个a + 无限个b ?

accepts infinite words that have finitely many symbols a



Büchi Automata

LTL Model Checking



# Model-Checking using Automata (1)

Transition systems can be transformed into Büchi automata, e.g.



## Model-Checking using Automata (2)

- The **states** of the automaton correspond to states of the transition system, plus one additional initial state.
- There is one transition from the new (and only!) initial state to all the states corresponding to initial states in the transition system.
- All other transitions of the automaton correspond to transitions in the transition system.
- The **alphabet** of the automaton is the set of all subsets of *Prop*, that is  $\mathcal{P}(\text{Prop})$ .
- The labels on automaton transitions are inherited from the *target states* in the transition system.
- All automaton states are accepting. (So all runs will be accepting!)

## Model-Checking using Automata (3)

**Question:** What are the words accepted by the resulting automaton?

**Answer:** Exactly those infinite sequences of sets of atomic propositions that occur along computation paths through the transition system !

For example:

- atomic propositions:  $Prop = \{n_1, n_2, t_1, t_2, c_1, c_2\}$
- $\{n_1, n_2\} \rightarrow \{t_1, n_2\} \rightarrow \{c_1, n_2\} \rightarrow \{n_1, n_2\} \rightarrow \dots$  occurs along a path through the transition system
- $\{n_1, n_2\}, \{t_1, n_2\}, \{c_1, n_2\}, \{n_1, n_2\}, \dots$  is accepted by the automaton

**Intuition:** runs of the automaton correspond to paths through the transition system.

## Model-Checking using Automata (3)

**Question:** What are the words accepted by the resulting automaton?

**Answer:** Exactly those infinite sequences of sets of atomic propositions that occur along computation paths through the transition system !

For example:

- atomic propositions:  $Prop = \{n_1, n_2, t_1, t_2, c_1, c_2\}$
- $\{n_1, n_2\} \rightarrow \{t_1, n_2\} \rightarrow \{c_1, n_2\} \rightarrow \{n_1, n_2\} \rightarrow \dots$  occurs along a path through the transition system
- $\{n_1, n_2\}, \{t_1, n_2\}, \{c_1, n_2\}, \{n_1, n_2\}, \dots$  is accepted by the automaton

**Intuition:** runs of the automaton correspond to paths through the transition system.

## Model-Checking using Automata (3)

**Question:** What are the words accepted by the resulting automaton?

**Answer:** Exactly those infinite sequences of sets of atomic propositions that occur along computation paths through the transition system !

For example:

- atomic propositions:  $Prop = \{n_1, n_2, t_1, t_2, c_1, c_2\}$
- $\{n_1, n_2\} \rightarrow \{t_1, n_2\} \rightarrow \{c_1, n_2\} \rightarrow \{n_1, n_2\} \rightarrow \dots$  occurs along a path through the transition system
- $\{n_1, n_2\}, \{t_1, n_2\}, \{c_1, n_2\}, \{n_1, n_2\}, \dots$  is accepted by the automaton

**Intuition:** runs of the automaton correspond to paths through the transition system.

## Model-Checking using Automata (4)

- LTL formula involving *Prop* can *also* be transformed into an automaton over the alphabet  $\mathcal{P}(\textit{Prop})$ .

Intuition:

- infinite words over  $\mathcal{P}(\textit{Prop})$  describe sequences of sets of atomic propositions
  - the resulting automaton should accept an infinite word precisely when the corresponding sequence satisfies the given formula
- 
- since the *model* and the *specification* are both automata, we can use the theory of Büchi automata to do model-checking !

## Model-Checking using Automata (4)

- LTL formula involving *Prop* can *also* be transformed into an automaton over the alphabet  $\mathcal{P}(\textit{Prop})$ .

Intuition:

- infinite words over  $\mathcal{P}(\textit{Prop})$  describe sequences of sets of atomic propositions
- the resulting automaton should accept an infinite word precisely when the corresponding sequence satisfies the given formula
- since the **model** and the **specification** are both automata, we can use the theory of Büchi automata to do model-checking !


# From LTL to Automata (Mutual Exclusion Property)

**A G**  $\neg(c_1 \wedge c_2)$  becomes



where:

- $\bigcirc \xrightarrow{c_1 \wedge c_2} \bigcirc$  stands for all arcs labelled by both  $c_1$  and  $c_2$ :



- $\bigcirc \xrightarrow{\epsilon} \bigcirc$  stands for *all* arcs labelled with subsets of *Prop*
- $\bigcirc \xrightarrow{\neg(c_1 \wedge c_2)} \bigcirc$  stands for *all* arcs labelled with subsets of *Prop* not containing both  $c_1$  and  $c_2$




# From LTL to Automata (Mutual Exclusion Property)

$\mathbf{A\ G} \neg(c_1 \wedge c_2)$  becomes



where:

- $\bigcirc \xrightarrow{c_1 \wedge c_2} \bigcirc$  stands for all arcs labelled by both  $c_1$  and  $c_2$ :



- $\bigcirc \xrightarrow{tt} \bigcirc$  stands for *all* arcs labelled with subsets of *Prop*
- $\bigcirc \xrightarrow{\neg(c_1 \wedge c_2)} \bigcirc$  stands for *all* arcs labelled with subsets of *Prop* not containing both  $c_1$  and  $c_2$

# From LTL to Automata

- mutual exclusion property: **A G**  $\neg(c_1 \wedge c_2)$



This automaton accepts all infinite words not containing both  $c_1$  and  $c_2$  at the same time (i.e. within the same "symbol").

- liveness property: **A F**  $c_1$ :



This automaton accepts all infinite words eventually containing  $c_1$  as part of a "symbol".

# From LTL to Automata

- mutual exclusion property:  $\mathbf{A\ G} \neg(c_1 \wedge c_2)$



This automaton accepts all infinite words not containing both  $c_1$  and  $c_2$  at the same time (i.e. within the same "symbol").

- liveness property:  $\mathbf{A\ F} c_1$ :

other examples



This automaton accepts all infinite words eventually containing  $c_1$  as part of a "symbol".

## Model Checking Using Automata (5)

- model captured by Büchi automaton  $\mathbb{A}$
- specification (LTL formula) captured by Büchi automaton  $\mathbb{S}$

- **Key observation:**

model satisfies the specification iff  $L(\mathbb{A}) \subseteq L(\mathbb{S})$  !

(any behaviour in the model satisfies the specification)

- thus, checking property " $\mathbb{S}$ " on model " $\mathbb{A}$ " reduces to checking the language inclusion  $L(\mathbb{A}) \subseteq L(\mathbb{S})$
- $L(\mathbb{A}) \subseteq L(\mathbb{S})$  can, in turn, be checked by checking:

$$L(\mathbb{A}) \cap \overline{L(\mathbb{S})} = \emptyset$$

or equivalently:

$$L(\mathbb{A} \cap \overline{\mathbb{S}}) = \emptyset$$

## Model Checking Using Automata (5)

- model captured by Büchi automaton  $\mathbb{A}$
- specification (LTL formula) captured by Büchi automaton  $\mathbb{S}$

- **Key observation:**

model satisfies the specification iff  $L(\mathbb{A}) \subseteq L(\mathbb{S})$  !

(any behaviour in the model satisfies the specification)

- thus, checking property " $\mathbb{S}$ " on model " $\mathbb{A}$ " reduces to checking the language inclusion  $L(\mathbb{A}) \subseteq L(\mathbb{S})$
- $L(\mathbb{A}) \subseteq L(\mathbb{S})$  can, in turn, be checked by checking:

$$L(\mathbb{A}) \cap \overline{L(\mathbb{S})} = \emptyset$$

or equivalently:

$$L(\mathbb{A} \cap \overline{\mathbb{S}}) = \emptyset$$

## Model Checking Using Automata (5)

- model captured by Büchi automaton  $\mathbb{A}$
- specification (LTL formula) captured by Büchi automaton  $\mathbb{S}$

- **Key observation:**

model satisfies the specification iff  $L(\mathbb{A}) \subseteq L(\mathbb{S})$  !

(any behaviour in the model satisfies the specification)

- thus, checking property " $\mathbb{S}$ " on model " $\mathbb{A}$ " reduces to checking the language inclusion  $L(\mathbb{A}) \subseteq L(\mathbb{S})$
- $L(\mathbb{A}) \subseteq L(\mathbb{S})$  can, in turn, be checked by checking:

$$L(\mathbb{A}) \cap \overline{L(\mathbb{S})} = \emptyset$$

or equivalently:

$$L(\mathbb{A} \cap \overline{\mathbb{S}}) = \emptyset$$

## Model Checking Using Automata (5)

- model captured by Büchi automaton  $\mathbb{A}$
- specification (LTL formula) captured by Büchi automaton  $\mathbb{S}$

- **Key observation:**

model satisfies the specification iff  $L(\mathbb{A}) \subseteq L(\mathbb{S})$  !

(any behaviour in the model satisfies the specification)

- thus, checking property " $\mathbb{S}$ " on model " $\mathbb{A}$ " reduces to checking the language inclusion  $L(\mathbb{A}) \subseteq L(\mathbb{S})$
- $L(\mathbb{A}) \subseteq L(\mathbb{S})$  can, in turn, be checked by checking:

$$L(\mathbb{A}) \cap \overline{L(\mathbb{S})} = \emptyset$$

or equivalently:

$$L(\mathbb{A} \cap \overline{\mathbb{S}}) = \emptyset$$

## Model Checking Using Automata (6)

Need to compute automaton  $\mathcal{A} \cap \overline{\mathcal{S}} \dots$

- computing  $\mathcal{A} \cap \mathcal{B}$  is relatively easy (polynomial complexity) ...
- ... but the complexity of computing  $\overline{\mathcal{S}}$  is exponential in the number of states if  $\mathcal{S}$  is *non-deterministic* !
  - what if  $\mathcal{S}$  is *deterministic* ?
  - **Note:** not any non-deterministic Büchi automaton has an equivalent deterministic one !
- better to directly generate automaton for the **negation** of the LTL property to check !



## Model Checking Using Automata (6)

Need to compute automaton  $\mathbb{A} \cap \overline{\mathbb{S}} \dots$

- computing  $\mathbb{A} \cap \mathbb{B}$  is relatively easy (polynomial complexity) ...
- ... but the complexity of computing  $\overline{\mathbb{S}}$  is exponential in the number of states if  $\mathbb{S}$  is *non-deterministic* !
  - what if  $\mathbb{S}$  is *deterministic* ?
  - **Note:** not any non-deterministic Büchi automaton has an equivalent deterministic one !
- better to directly generate automaton for the *negation* of the LTL property to check !

## Model Checking Using Automata (6)

Need to compute automaton  $\mathcal{A} \cap \overline{\mathcal{S}} \dots$

- computing  $\mathcal{A} \cap \mathcal{B}$  is relatively easy (polynomial complexity) ...
- ... but the complexity of computing  $\overline{\mathcal{S}}$  is exponential in the number of states if  $\mathcal{S}$  is *non-deterministic* !
  - what if  $\mathcal{S}$  is *deterministic* ?
  - **Note:** not any non-deterministic Büchi automaton has an equivalent deterministic one !
- better to directly generate automaton for the *negation* of the LTL property to check !

## Model Checking Using Automata (6)

Need to compute automaton  $A \cap \bar{S} \dots$


- computing  $A \cap B$  is relatively easy (polynomial complexity) ...
- ... but the complexity of computing  $\bar{S}$  is exponential in the number of states if  $S$  is *non-deterministic* !
  - what if  $S$  is *deterministic* ?
  - **Note:** not any non-deterministic Büchi automaton has an equivalent deterministic one !
- better to directly generate automaton for the **negation** of the LTL property to check !

# Example 1

- automaton for negation of mutual exclusion property:



- automaton for model of mutual exclusion:




# Example 1

- automaton for negation of mutual exclusion property:



combine



- automaton for model of mutual exclusion:



# Construction of the Product Automaton

- the states of the product automaton are given by *pairs of states* of the two automata,


**Note:** this is only possible if all the states of one of the automata are accepting.

-  precisely when  and 
- $(s_1, s_2)$  is an initial state precisely when both  $s_1$  and  $s_2$  are initial states,
- $(s_1, s_2)$  is an accepting state precisely when both  $s_1$  and  $s_2$  are accepting states.

# Construction of the Product Automaton

- the states of the product automaton are given by *pairs of states* of the two automata,

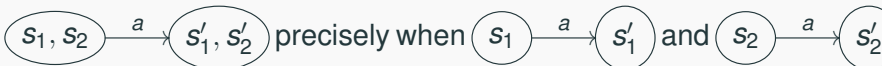

**Note:** this is only possible if all the states of one of the automata are accepting.

-  precisely when  $s_1 \xrightarrow{a} s'_1$  and  $s_2 \xrightarrow{a} s'_2$
- $(s_1, s_2)$  is an initial state precisely when both  $s_1$  and  $s_2$  are initial states,
- $(s_1, s_2)$  is an accepting state precisely when both  $s_1$  and  $s_2$  are accepting states.

# Construction of the Product Automaton

- the states of the product automaton are given by *pairs of states* of the two automata,

**Note:** this is only possible if all the states of one of the automata are accepting.

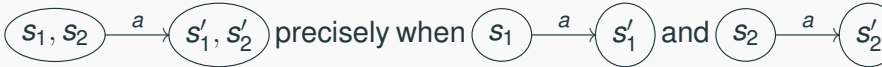
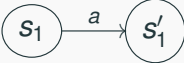
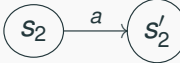
- precisely when 
- $(s_1, s_2)$  is an initial state precisely when both  $s_1$  and  $s_2$  are initial states,
- $(s_1, s_2)$  is an accepting state precisely when both  $s_1$  and  $s_2$  are accepting states.



# Construction of the Product Automaton


- the states of the product automaton are given by *pairs of states* of the two automata,

**Note:** this is only possible if all the states of one of the automata are accepting.

- precisely when  and 
- $(s_1, s_2)$  is an initial state precisely when both  $s_1$  and  $s_2$  are initial states,
- $(s_1, s_2)$  is an accepting state precisely when both  $s_1$  and  $s_2$  are accepting states.

## Example 1 (Cont'd)


Product automaton has no accepting states, hence does not accept *any* infinite word:



Since the language accepted by the product automaton is **empty**, it follows that the mutual exclusion property  $\mathbf{AG} \neg(c_1 \wedge c_2)$  **holds** in the original transition system.

## Example 1 (Cont'd)

Product automaton has no accepting states, hence does not accept *any* infinite word:



Since the language accepted by the product automaton is **empty**, it follows that the mutual exclusion property **A G  $\neg(c_1 \wedge c_2)$**  **holds** in the original transition system.

## Example 2

- automaton for negation of liveness property:




Note: this is not deterministic!

c1 never happen

buchi automaton is used for

infinte path

- automaton for model of mutual exclusion:




## Example 2

- automaton for negation of liveness property:




Note: this is not  
**deterministic!**

- automaton for model of mutual exclusion:



## Example 2 (Cont'd)

Product automaton:




Since the language accepted by the product automaton is **not empty** (why?), it follows that the liveness property  **$A \text{ } F c_1$  does not hold** in the original transition system.

Any infinite word accepted by the above automaton produces a **counterexample!**

## Example 2 (Cont'd)

Product automaton:




Since the language accepted by the product automaton is **not empty** (why?), it follows that the liveness property  **$A \text{ } \mathbf{F} \text{ } c_1$**  **does not hold** in the original transition system.

Any infinite word accepted by the above automaton produces a **counterexample!**

## Example 2 (Cont'd)

Product automaton:



Since the language accepted by the product automaton is **not empty** (why?), it follows that the liveness property  **$A \ F \ c_1$**  **does not hold** in the original transition system.

Any infinite word accepted by the above automaton produces a **counterexample!**



# LTL Model-Checking Algorithm

- there is an algorithm for translating LTL formulas into Büchi automata (details in [Clarke et al.]),
- the complexity of the resulting model-checking algorithm is **linear** in the size of the model (and **exponential** in the length of the formula),
- fairness assumptions can also be described by LTL formulas.

### 命题的否定

do not enter c infinite time

1. Construct an LTL formula and a Büchi automaton for the correctness property: *"In each execution, each process enters the critical section an infinite number of times."*
2. Check, by manually applying the model checking algorithm, whether this property holds in the model on slide 10.

# The SPIN Model Checker

- example of **explicit-state** model checker
- uses **Promela** modelling language to describe (concurrent/distributed) software systems
- uses LTL as main mechanism for specifying correctness . . .
- . . . but annotations to Promela models also used for simple safety/liveness properties
- can handle systems with millions of states

# Benefits and Limitations of Explicit State Model Checking

## Benefits:

- automatic
- exhaustive
- produces counter-examples

## Limitations:

- state explosion problem – partially addressed by **symbolic model checking** (our next topic)
- only works for finite-state systems