

COMP6210 Automated Software Verification

Automata-Based Verification for LTL

Pavel Naumov

Intended Learning Outcomes

By the end of this lecture, you will be able to

- define a Büchi automaton
- construct Buchi automata for simple LTL formulas
- explain how Buchi automata are used in model-checking LTL formulas


Büchi Automata

LTL Model Checking

Recap on Finite Automata (1)

different btw fa and ts

- **finite automata** accept/reject *strings* over a given *alphabet*
- equivalent to regular expressions
- a finite automaton over the alphabet $\{a, b\}$:


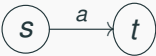




- the initial state is s_1
- the final state is s_1
- the string **abba** is accepted
- the string **bab** is rejected

Recap on Finite Automata (2)

Let Σ be a finite **alphabet** (consisting symbols that we will use in strings/words over Σ).

A **finite automaton** \mathbb{A} consists of:

- a finite set of **states**: 
- a finite set of **transitions** between states, labelled by symbols from Σ : 
- a subset of **initial states**: 
- a subset of **final/accepting states**: 

Recap on Finite Automata (3)

- A **run** of an automaton on a finite word $a_1 \dots a_n$ is a finite path through the automaton starting in an initial state, with transitions labelled by a_1, \dots, a_n (in this order):

$$s_0 \xrightarrow{a_1} s_1 \xrightarrow{a_2} \dots \xrightarrow{a_{n-1}} s_{n-1} \xrightarrow{a_n} s_n$$

- An automaton \mathbb{A} **accepts** a finite word $a_1 \dots a_n$ if there exists a run of \mathbb{A} on a_1, \dots, a_n which ends in an accepting state.
- The **language** of an automaton \mathbb{A} , denoted $L(\mathbb{A})$, consists of all the words it accepts.


Recap on Finite Automata (3)

- A **run** of an automaton on a finite word $a_1 \dots a_n$ is a finite path through the automaton starting in an initial state, with transitions labelled by a_1, \dots, a_n (in this order):

$$s_0 \xrightarrow{a_1} s_1 \xrightarrow{a_2} \dots \xrightarrow{a_{n-1}} s_{n-1} \xrightarrow{a_n} s_n$$

- An automaton \mathbb{A} **accepts** a finite word $a_1 \dots a_n$ if there **exists** a run of \mathbb{A} on a_1, \dots, a_n which ends in an accepting state.
- The **language** of an automaton \mathbb{A} , denoted $L(\mathbb{A})$, consists of all the words it accepts.

Recap on Finite Automata (4)



- some runs of this automaton:

$$s_1 \xrightarrow{b} s_2 \xrightarrow{b} s_2 \xrightarrow{a} s_1$$


$$s_1 \xrightarrow{a} s_1 \xrightarrow{b} s_2 \xrightarrow{b} s_2 \xrightarrow{a} s_1$$

- some words accepted by this automaton:

ϵ (the empty word), a , ba , bba , $baba$, ...

- the language of this automaton consists of:
 - the empty word,
 - all non-empty words that end with an a .

Finite Automata over Infinite Words (Büchi Automata)




(The initial state is s_1 .)

- An **infinite run** is an infinite path through the automaton, starting in an initial state.

- e.g. $s_0 \xrightarrow{a_1} s_1 \xrightarrow{a_2} s_2 \xrightarrow{a_3} \dots$

- An automaton **accepts** an infinite word $a_1 a_2 a_3 \dots$ if there exists a run of the automaton labelled by a_1, a_2, \dots , which passes through an accepting state **infinitely often**. We call such a run **accepting**.

Finite Automata over Infinite Words (Büchi Automata)




(The initial state is s_1 .)

- An **infinite run** is an infinite path through the automaton, starting in an initial state.

- e.g. $s_0 \xrightarrow{a_1} s_1 \xrightarrow{a_2} s_2 \xrightarrow{a_3} \dots$

- An automaton **accepts** an infinite word $a_1 a_2 a_3 \dots$ if there exists a run of the automaton labelled by a_1, a_2, \dots , which passes through an accepting state **infinitely often**. We call such a run **accepting**.

Example




(The only initial state and only final state is s_1 .)

The run $s_1 \xrightarrow{b} s_2 \xrightarrow{a} s_1 \xrightarrow{b} s_2 \xrightarrow{a} \dots$ is accepting.

The run $s_1 \xrightarrow{a} s_1 \xrightarrow{b} s_2 \xrightarrow{b} s_2 \xrightarrow{b} \dots$ is not accepting.

The language of this Büchi automaton consists of all **infinite** words containing infinitely many a symbols.

Example



(The only initial state and only final state is s_1 .)

The run $s_1 \xrightarrow{b} s_2 \xrightarrow{a} s_1 \xrightarrow{b} s_2 \xrightarrow{a} \dots$ is accepting.

The run $s_1 \xrightarrow{a} s_1 \xrightarrow{b} s_2 \xrightarrow{b} s_2 \xrightarrow{b} \dots$ is not accepting.

The language of this Büchi automaton consists of all **infinite** words containing infinitely many a symbols.

Some Properties of Büchi Automata

If \mathbb{A} and \mathbb{B} are automata, then there exist:


- an automaton $\mathbb{A} \cap \mathbb{B}$ (the **product** automaton) such that
$$L(\mathbb{A} \cap \mathbb{B}) = L(\mathbb{A}) \cap L(\mathbb{B}).$$
- an automaton $\overline{\mathbb{A}}$ (the **complement** automaton) such that $L(\overline{\mathbb{A}})$ contains exactly those infinite words which are **not** in $L(\mathbb{A})$.

Some Properties of Büchi Automata

If \mathbb{A} and \mathbb{B} are automata, then there exist:


- an automaton $\mathbb{A} \cap \mathbb{B}$ (the **product** automaton) such that
$$L(\mathbb{A} \cap \mathbb{B}) = L(\mathbb{A}) \cap L(\mathbb{B}).$$
- an automaton $\overline{\mathbb{A}}$ (the **complement** automaton) such that $L(\overline{\mathbb{A}})$ contains exactly those infinite words which are **not** in $L(\mathbb{A})$.

Buchi Automaton Complement Construction



accepts infinite words that have infinitely many symbols a

how to change that ?



无限个a + 无限个b ?



accepts infinite words that have finitely many symbols a

Büchi Automata

LTL Model Checking

Model-Checking using Automata (1)

Transition systems can be transformed into Büchi automata, e.g.



Model-Checking using Automata (2)

- The **states** of the automaton correspond to states of the transition system, plus one additional initial state.
- There is one transition from the new (and only!) initial state to all the states corresponding to initial states in the transition system.
- All other transitions of the automaton correspond to transitions in the transition system.
- The **alphabet** of the automaton is the set of all subsets of *Prop*, that is $\mathcal{P}(\text{Prop})$.
- The labels on automaton transitions are inherited from the *target states* in the transition system.
- All automaton states are accepting. (So all runs will be accepting!)

Model-Checking using Automata (3)

Question: What are the words accepted by the resulting automaton?

Answer: Exactly those infinite sequences of sets of atomic propositions that occur along computation paths through the transition system !

For example:

- atomic propositions: $Prop = \{n_1, n_2, t_1, t_2, c_1, c_2\}$
- $\{n_1, n_2\} \rightarrow \{t_1, n_2\} \rightarrow \{c_1, n_2\} \rightarrow \{n_1, n_2\} \rightarrow \dots$ occurs along a path through the transition system
- $\{n_1, n_2\}, \{t_1, n_2\}, \{c_1, n_2\}, \{n_1, n_2\}, \dots$ is accepted by the automaton

Intuition: runs of the automaton correspond to paths through the transition system.

Model-Checking using Automata (3)

Question: What are the words accepted by the resulting automaton?

Answer: Exactly those infinite sequences of sets of atomic propositions that occur along computation paths through the transition system !

For example:

- atomic propositions: $Prop = \{n_1, n_2, t_1, t_2, c_1, c_2\}$
- $\{n_1, n_2\} \rightarrow \{t_1, n_2\} \rightarrow \{c_1, n_2\} \rightarrow \{n_1, n_2\} \rightarrow \dots$ occurs along a path through the transition system
- $\{n_1, n_2\}, \{t_1, n_2\}, \{c_1, n_2\}, \{n_1, n_2\}, \dots$ is accepted by the automaton

Intuition: runs of the automaton correspond to paths through the transition system.

Model-Checking using Automata (3)

Question: What are the words accepted by the resulting automaton?

Answer: Exactly those infinite sequences of sets of atomic propositions that occur along computation paths through the transition system !

For example:

- atomic propositions: $Prop = \{n_1, n_2, t_1, t_2, c_1, c_2\}$
- $\{n_1, n_2\} \rightarrow \{t_1, n_2\} \rightarrow \{c_1, n_2\} \rightarrow \{n_1, n_2\} \rightarrow \dots$ occurs along a path through the transition system
- $\{n_1, n_2\}, \{t_1, n_2\}, \{c_1, n_2\}, \{n_1, n_2\}, \dots$ is accepted by the automaton

Intuition: runs of the automaton correspond to paths through the transition system.

Model-Checking using Automata (4)

- LTL formula involving *Prop* can *also* be transformed into an automaton over the alphabet $\mathcal{P}(\textit{Prop})$.

Intuition:

- infinite words over $\mathcal{P}(\textit{Prop})$ describe sequences of sets of atomic propositions
- the resulting automaton should accept an infinite word precisely when the corresponding sequence satisfies the given formula
- since the *model* and the *specification* are both automata, we can use the theory of Büchi automata to do model-checking !

Model-Checking using Automata (4)

- LTL formula involving *Prop* can *also* be transformed into an automaton over the alphabet $\mathcal{P}(\textit{Prop})$.

Intuition:

- infinite words over $\mathcal{P}(\textit{Prop})$ describe sequences of sets of atomic propositions
- the resulting automaton should accept an infinite word precisely when the corresponding sequence satisfies the given formula
- since the **model** and the **specification** are both automata, we can use the theory of Büchi automata to do model-checking !


From LTL to Automata (Mutual Exclusion Property)

A G $\neg(c_1 \wedge c_2)$ becomes



where:

- $\bigcirc \xrightarrow{c_1 \wedge c_2} \bigcirc$ stands for all arcs labelled by both c_1 and c_2 :



- $\bigcirc \xrightarrow{\epsilon} \bigcirc$ stands for *all* arcs labelled with subsets of *Prop*
- $\bigcirc \xrightarrow{\neg(c_1 \wedge c_2)} \bigcirc$ stands for *all* arcs labelled with subsets of *Prop* not containing both c_1 and c_2


From LTL to Automata (Mutual Exclusion Property)

$\mathbf{A\ G} \neg(c_1 \wedge c_2)$ becomes



where:

- $\bigcirc \xrightarrow{c_1 \wedge c_2} \bigcirc$ stands for all arcs labelled by both c_1 and c_2 :



- $\bigcirc \xrightarrow{tt} \bigcirc$ stands for *all* arcs labelled with subsets of *Prop*
- $\bigcirc \xrightarrow{\neg(c_1 \wedge c_2)} \bigcirc$ stands for *all* arcs labelled with subsets of *Prop* not containing both c_1 and c_2

From LTL to Automata

- mutual exclusion property: **A G** $\neg(c_1 \wedge c_2)$



This automaton accepts all infinite words not containing both c_1 and c_2 at the same time (i.e. within the same "symbol").

- liveness property: **A F** c_1 :



This automaton accepts all infinite words eventually containing c_1 as part of a "symbol".

From LTL to Automata

- mutual exclusion property: $\mathbf{A\ G} \neg(c_1 \wedge c_2)$



This automaton accepts all infinite words not containing both c_1 and c_2 at the same time (i.e. within the same "symbol").

- liveness property: $\mathbf{A\ F} c_1$:

other examples



This automaton accepts all infinite words eventually containing c_1 as part of a "symbol".

Model Checking Using Automata (5)

- model captured by Büchi automaton \mathbb{A}
- specification (LTL formula) captured by Büchi automaton \mathbb{S}

- **Key observation:**

model satisfies the specification iff $L(\mathbb{A}) \subseteq L(\mathbb{S})$!

(any behaviour in the model satisfies the specification)

- thus, checking property " \mathbb{S} " on model " \mathbb{A} " reduces to checking the language inclusion $L(\mathbb{A}) \subseteq L(\mathbb{S})$
- $L(\mathbb{A}) \subseteq L(\mathbb{S})$ can, in turn, be checked by checking:

$$L(\mathbb{A}) \cap \overline{L(\mathbb{S})} = \emptyset$$

or equivalently:

$$L(\mathbb{A} \cap \overline{\mathbb{S}}) = \emptyset$$

Model Checking Using Automata (5)

- model captured by Büchi automaton \mathbb{A}
- specification (LTL formula) captured by Büchi automaton \mathbb{S}

- **Key observation:**

model satisfies the specification iff $L(\mathbb{A}) \subseteq L(\mathbb{S})$!

(any behaviour in the model satisfies the specification)

- thus, checking property " \mathbb{S} " on model " \mathbb{A} " reduces to checking the language inclusion $L(\mathbb{A}) \subseteq L(\mathbb{S})$
- $L(\mathbb{A}) \subseteq L(\mathbb{S})$ can, in turn, be checked by checking:

$$L(\mathbb{A}) \cap \overline{L(\mathbb{S})} = \emptyset$$

or equivalently:

$$L(\mathbb{A} \cap \overline{\mathbb{S}}) = \emptyset$$

Model Checking Using Automata (5)

- model captured by Büchi automaton \mathbb{A}
- specification (LTL formula) captured by Büchi automaton \mathbb{S}

- **Key observation:**

model satisfies the specification iff $L(\mathbb{A}) \subseteq L(\mathbb{S})$!

(any behaviour in the model satisfies the specification)

- thus, checking property " \mathbb{S} " on model " \mathbb{A} " reduces to checking the language inclusion $L(\mathbb{A}) \subseteq L(\mathbb{S})$
- $L(\mathbb{A}) \subseteq L(\mathbb{S})$ can, in turn, be checked by checking:

$$L(\mathbb{A}) \cap \overline{L(\mathbb{S})} = \emptyset$$

or equivalently:

$$L(\mathbb{A} \cap \overline{\mathbb{S}}) = \emptyset$$

Model Checking Using Automata (5)

- model captured by Büchi automaton \mathbb{A}
- specification (LTL formula) captured by Büchi automaton \mathbb{S}

- **Key observation:**

model satisfies the specification iff $L(\mathbb{A}) \subseteq L(\mathbb{S})$!

(any behaviour in the model satisfies the specification)

- thus, checking property " \mathbb{S} " on model " \mathbb{A} " reduces to checking the language inclusion $L(\mathbb{A}) \subseteq L(\mathbb{S})$
- $L(\mathbb{A}) \subseteq L(\mathbb{S})$ can, in turn, be checked by checking:

$$L(\mathbb{A}) \cap \overline{L(\mathbb{S})} = \emptyset$$

or equivalently:

$$L(\mathbb{A} \cap \overline{\mathbb{S}}) = \emptyset$$

Model Checking Using Automata (6)

Need to compute automaton $\mathcal{A} \cap \overline{\mathcal{S}} \dots$

- computing $\mathcal{A} \cap \mathcal{B}$ is relatively easy (polynomial complexity) ...
- ... but the complexity of computing $\overline{\mathcal{S}}$ is exponential in the number of states if \mathcal{S} is *non-deterministic* !
 - what if \mathcal{S} is *deterministic* ?
 - **Note:** not any non-deterministic Büchi automaton has an equivalent deterministic one !
- better to directly generate automaton for the **negation** of the LTL property to check !

Model Checking Using Automata (6)

Need to compute automaton $\mathbb{A} \cap \overline{\mathbb{S}} \dots$

- computing $\mathbb{A} \cap \mathbb{B}$ is relatively easy (polynomial complexity) ...
- ... but the complexity of computing $\overline{\mathbb{S}}$ is exponential in the number of states if \mathbb{S} is *non-deterministic* !
 - what if \mathbb{S} is *deterministic* ?
 - **Note:** not any non-deterministic Büchi automaton has an equivalent deterministic one !
- better to directly generate automaton for the *negation* of the LTL property to check !

Model Checking Using Automata (6)

Need to compute automaton $A \cap \bar{S} \dots$

- computing $A \cap B$ is relatively easy (polynomial complexity) ...
- ... but the complexity of computing \bar{S} is exponential in the number of states if S is *non-deterministic* !
 - what if S is *deterministic* ?
 - **Note:** not any non-deterministic Büchi automaton has an equivalent deterministic one !
- better to directly generate automaton for the *negation* of the LTL property to check !

Model Checking Using Automata (6)

Need to compute automaton $\mathcal{A} \cap \overline{\mathcal{S}} \dots$


- computing $\mathcal{A} \cap \mathcal{B}$ is relatively easy (polynomial complexity) ...
- ... but the complexity of computing $\overline{\mathcal{S}}$ is exponential in the number of states if \mathcal{S} is *non-deterministic* !
 - what if \mathcal{S} is *deterministic* ?
 - **Note:** not any non-deterministic Büchi automaton has an equivalent deterministic one !
- better to directly generate automaton for the **negation** of the LTL property to check !

Example 1

- automaton for negation of mutual exclusion property:



- automaton for model of mutual exclusion:




Example 1

- automaton for negation of mutual exclusion property:



combine



- automaton for model of mutual exclusion:



Construction of the Product Automaton

- the states of the product automaton are given by *pairs of states* of the two automata,


Note: this is only possible if all the states of one of the automata are accepting.

-  precisely when  and 
- (s_1, s_2) is an initial state precisely when both s_1 and s_2 are initial states,
- (s_1, s_2) is an accepting state precisely when both s_1 and s_2 are accepting states.

Construction of the Product Automaton

- the states of the product automaton are given by *pairs of states* of the two automata,

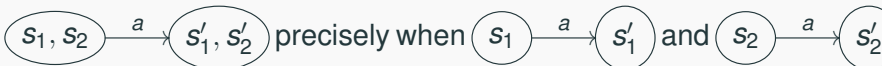

Note: this is only possible if all the states of one of the automata are accepting.

-  precisely when $s_1 \xrightarrow{a} s'_1$ and $s_2 \xrightarrow{a} s'_2$
- (s_1, s_2) is an initial state precisely when both s_1 and s_2 are initial states,
- (s_1, s_2) is an accepting state precisely when both s_1 and s_2 are accepting states.

Construction of the Product Automaton

- the states of the product automaton are given by *pairs of states* of the two automata,

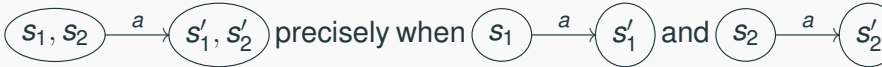
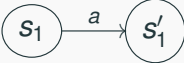
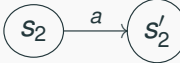
Note: this is only possible if all the states of one of the automata are accepting.

- precisely when 
- (s_1, s_2) is an initial state precisely when both s_1 and s_2 are initial states,
- (s_1, s_2) is an accepting state precisely when both s_1 and s_2 are accepting states.

Construction of the Product Automaton


- the states of the product automaton are given by *pairs of states* of the two automata,

Note: this is only possible if all the states of one of the automata are accepting.

- precisely when  and 
- (s_1, s_2) is an initial state precisely when both s_1 and s_2 are initial states,
- (s_1, s_2) is an accepting state precisely when both s_1 and s_2 are accepting states.

Example 1 (Cont'd)


Product automaton has no accepting states, hence does not accept *any* infinite word:



Since the language accepted by the product automaton is **empty**, it follows that the mutual exclusion property $\mathbf{A} \mathbf{G} \neg(c_1 \wedge c_2)$ **holds** in the original transition system.

Example 1 (Cont'd)

Product automaton has no accepting states, hence does not accept *any* infinite word:



Since the language accepted by the product automaton is **empty**, it follows that the mutual exclusion property **$\mathbf{A\ G} \neg(c_1 \wedge c_2)$** holds in the original transition system.

Example 2

- automaton for negation of liveness property:




Note: this is not deterministic!

c1 never happen

buchi automaton id used for infinte path

- automaton for model of mutual exclusion:




Example 2

- automaton for negation of liveness property:




Note: this is not
deterministic!

- automaton for model of mutual exclusion:



Example 2 (Cont'd)

Product automaton:




Since the language accepted by the product automaton is **not empty** (why?), it follows that the liveness property **$A \ F \ c_1$** **does not hold** in the original transition system.

Any infinite word accepted by the above automaton produces a **counterexample!**

Example 2 (Cont'd)

Product automaton:




Since the language accepted by the product automaton is **not empty** (why?), it follows that the liveness property **$A \text{ } \mathbf{F} \text{ } c_1$** **does not hold** in the original transition system.

Any infinite word accepted by the above automaton produces a counterexample!

Example 2 (Cont'd)

Product automaton:



Since the language accepted by the product automaton is **not empty** (why?), it follows that the liveness property **$A \ F \ c_1$** **does not hold** in the original transition system.

Any infinite word accepted by the above automaton produces a **counterexample!**

LTL Model-Checking Algorithm

- there is an algorithm for translating LTL formulas into Büchi automata (details in [Clarke et al.]),
- the complexity of the resulting model-checking algorithm is **linear** in the size of the model (and **exponential** in the length of the formula),
- fairness assumptions can also be described by LTL formulas.

Exercise

1. Construct an LTL formula and a Büchi automaton for the correctness property: *"In each execution, each process enters the critical section an infinite number of times."*
2. Check, by manually applying the model checking algorithm, whether this property holds in the model on slide 10.

The SPIN Model Checker

- example of **explicit-state** model checker
- uses **Promela** modelling language to describe (concurrent/distributed) software systems
- uses LTL as main mechanism for specifying correctness . . .
- . . . but annotations to Promela models also used for simple safety/liveness properties
- can handle systems with millions of states

Benefits and Limitations of Explicit State Model Checking

Benefits:

- automatic
- exhaustive
- produces counter-examples

Limitations:

- state explosion problem – partially addressed by **symbolic model checking** (our next topic)
- only works for finite-state systems